

WHEN DOES EXPERIMENTAL KNOWLEDGE IMPROVE SCIENTIFIC DECISIONS?

Anonymous authors

Paper under double-blind review

ABSTRACT

Scientific agents can improve experimental outcomes without converting experience into accurate, executable and useful knowledge. We measure four observable conversion questions in ChemWorld: search to selective correction, numerical prediction to structure, predictions to executable laws, and laws to unseen decisions. Two model systems each entered 135 campaigns across 45 matched task–world clusters. Prediction errors fell, selective-correction criteria remained unmet, and executable summaries lost predictive information. A subsequent 120-session disclosure intervention in ten shared worlds retained 112 valid completions and eight failures or interruptions. Completing the public observation mapping raised joint family/exponent recovery from 6/20 to 20/20 for GPT and 3/20 to 17/20 for DeepSeek: a pooled gain of 70 percentage points (approximate 95% interval [52.5, 85]). Optimal action choices increased from 5/20 to 12/20 for GPT but changed from 6/20 to 4/20 for DeepSeek. Structure recovery therefore did not ensure consistent action gains. Separately, a ten-world factorial intervention found no supported material fitted-law benefit and complete fitted-law agent/maximizer agreement. Delivering model laws alone to fresh recipients on new same-world candidates reduced regret by 0.13723 relative to task information alone, while nearest-evidence retrieval attained zero regret. These results identify conditional information benefits and distinguish artifact fidelity from decision utility, without establishing four internal causal failures, universal agent deficits or transfer to changed physical conditions.


1 INTRODUCTION

Autonomous scientific agents choose experiments, interpret observations and recommend what to do next. The value of experimental knowledge therefore depends on the decisions it supports. A useful outcome, an accurate prediction and an executable scientific summary answer different questions: an agent may find a productive recipe, predict a local response, or express a relation without being able to select a good plan under new conditions.

This problem matters as language-model agents operate chemistry tools and self-driving laboratories (Boiko et al., 2023; Bran et al., 2024; Szymanski et al., 2023; Darvish et al., 2025; Song et al., 2025; Vriza et al., 2026). Interactive discovery environments make repeated experimentation accessible (Jansen et al., 2024; Gandhi et al., 2025; Duan et al., 2025; Zheng et al., 2026; Yang et al., 2026; Batzoglou, 2026). Predict-then-optimize and decision-focused learning already establish that predictive error and downstream decision loss can differ (Elmachtoub & Grigas, 2022; Wilder et al., 2019). The additional question for a scientific agent is how autonomously acquired evidence, a supplied prior and a submitted knowledge artifact relate to the actual decision reached by the complete system.

We use ChemWorld to make these objects separately observable. Within a matched cluster, the external world, public operations and resources remain fixed while the supplied initial description is opaque, aligned or misspecified at a declared entity, parametric or structural locus (Qiu et al., 2026). A persistent session performs experiments and submits predictions and typed laws. Independent evaluators score those artifacts and execute complete action plans. The assignment changes participant-facing information; it does not directly manipulate an unobservable internal belief. The foundation paper establishes the bounded environment and replay semantics, while this study evaluates complete agent–tool configurations.

From experimental evidence to scientific decisions



Four separate measurements; success in one does not establish success in another.

Figure 1: **Endpoint success does not reveal what the agent learned.** ChemWorld varies the supplied initial description while fixing the executable world, public rules, and budget. Four conversion questions separately assess correction, structure, law fidelity and unseen decisions. The diagram describes the design; its arrows do not identify internal belief or causal mediation through submitted knowledge.

We organize the study around four observable conversion questions: **F1**, experimental search to selective correction; **F2**, numerical prediction to structural identification; **F3**, conditional predictions to executable laws; and **F4**, submitted laws to unseen-plan decisions. Each question uses its own protocol, denominator and failure accounting. The four questions do not form an identified internal causal chain, and success at one readout does not substitute for measuring the next.

The primary study used a fixed DeepSeek-v4-flash experimental-agent configuration, with 135 scheduled cells nested within 45 task-world clusters; a GPT-5.6-sol successor used the same scientific surface. These are within-system intervention studies in two model families, not a representative model survey. A matched-packet diagnostic reduces submission demands; a subsequent disclosure intervention tests whether completing the observation mapping improves structure and action. We also cross law source with decision rule under fixed evidence, then deliver laws independently of raw evidence to fresh recipients. These interventions exhibit settings with agreement and useful knowledge, bounding a general law-use failure account without identifying the causes of historical cross-protocol differences.

2 ENVIRONMENT AND EVALUATION

2.1 CHEMWORLD AS A CONTROLLED CAUSAL PROBE

ChemWorld separates three objects. The executable world contains hidden state, transition and measurement laws, resources, and terminal assays. The public task exposes typed operations, instruments, observations, and an initial model. The evaluator owns held-out truth, replay, and scoring. For world w and initial-model arm a , the hidden transition process remains fixed while the supplied model $M_{0,a}$ changes. The aligned model is correct at one declared locus, the misspecified model is explicit but wrong at that locus, and the opaque arm withholds the corresponding structure. The design changes participant-facing information while holding world physics and the available operation and measurement semantics fixed.

Actions are transactions, not free-form prose. The host validates a typed request, checks resources and preconditions, commits the state transition, and returns a public observation or a structured rejection. Failed actions and discards remain in the trajectory. Every submitted action, commit, rollback, measurement, assay, and resource delta is recorded for exact replay. These semantics matter: an action recommendation is scientific evidence only when the evaluator can execute the same complete plan without filling in hidden workflow choices.

2.2 OBSERVABLE OUTCOMES AND MEASUREMENT LIMITS

We distinguish five downstream outcomes. **Evidence acquisition** asks which informative experiments are selected and observed. **Numerical revision** asks whether held-out counterfactual prediction error falls. **Structural identification** asks whether the agent rejects a false form and recovers the registered relation. **Executable compression** asks whether a typed law preserves the information in conditional predictions. **Decision transfer** asks whether the final state selects a good previously unseen plan. A sixth requirement, **evaluator validity**, asks whether the control and metric actually measure the decision estimand. Success at one stage is neither defined as nor assumed to imply success at the next.

At five checkpoints per session, the participant reports initial-model reliability, predictions, uncertainty, evidence references, an executable law, and the next experimental intent. The evaluator executes prespecified counterfactual query sets independently of the participant. The DeepSeek-v4-flash surface scores 675/675 checkpoints from 420/420 truth executions; the GPT-5.6-sol surface scores 669/675 checkpoints from its own 420/420 truth executions. The evaluator later runs each available final typed law on the same coordinates. Paired blind replay evaluates the final recommendation against the observed incumbent, while a separate longitudinal assay reveals eight outcome-hidden complete ActionPlans only after autonomous exploration has ended.

3 EXPERIMENTAL PROGRAMME

The prospective programme is layer-stratified. Entity interventions cover five task families and five independently selected public worlds per task. Parametric and structural blocks each cover two validated task families and five worlds. Every task–world cluster contains opaque, aligned, and misspecified arms, yielding 135 separate sessions nested within 45 independent task–world clusters. Campaign length is locus-specific: eight, ten, or twelve complete experiments, with five checkpoints in every session. Exploratory, validation, prospective, matched-evidence, and open-action worlds remain separated. We abbreviate the prospective cohort as C2, the matched partition-packet diagnostic as B2, its typed-law/action control as B3, and the entity/parametric/structural loci as A-E/A-P/A-S.

The primary contrast tests selective evidence-driven correction. If $E_{a,k}^{(\ell)}$ is held-out error for arm a , checkpoint k , and locus ℓ , then

$$C_\ell = \left(E_{\text{mis},0}^{(\ell)} - E_{\text{mis},K}^{(\ell)} \right) - \left(E_{\text{aligned},0}^{(\ell)} - E_{\text{aligned},K}^{(\ell)} \right).$$

Success requires greater correction in the misspecified arm, improvement of that arm, and no material deterioration of the aligned arm. Loci are decided separately; unlike intervention semantics are not pooled. Failed cells stay in the scheduled denominator; confirmatory correction gates use adverse bounds, while last-observation and zero-improvement imputations are sensitivities. Only infrastructure failures without a persisted trajectory can resume under a fixed attempt cap.

Matched-evidence sessions use cloned worlds and give every arm the same counterevidence after a pre-response. They reveal conditional post-packet updating but, without a turn-matched no-packet control, do not identify a pure evidence-packet effect. The longitudinal action matrix separately contains three tasks, five worlds, and three arms (45 scheduled cells). After 12 autonomous experiments, each agent ranks eight new plans; regret and Top-1 are primary action readouts, while complete-rank correlation and law adequacy are diagnostics. Table 1 keeps these layers and their claim boundaries explicit.

C2 and B3 cover both configurations; the original longitudinal action assay is DeepSeek-only. Later artifact and disclosure interventions test both systems under block-specific reasoning settings. Model panels show within-system contrasts, not equal-compute rankings. Prespecified model/repeat means within world estimate a fixed mixture of these systems; shared worlds are not independent replications across models. The appendix maps model coverage to each question.

Table 1: Executed evidence layers and claim boundaries. A completed work package may contain a retained scientific rejection; unstarted units are not silently removed.

Layer	Units	Execution	Supported role
Prospective C2	270 scheduled	DeepSeek 121 complete; GPT 126 complete	Search, prediction, law, incumbent replay
Matched evidence	75 sessions	DeepSeek high + GPT medium (60); DeepSeek low (15)	Conditional numerical–exact-law-expression dissociation
Historical B3	60 scheduled	DeepSeek 17+13 failures; GPT 30	Structural recovery and action bridge
Minimal-interface diagnostic	120 scheduled; 5 reused worlds	117 complete; 3 failures	Structure and optional-numeric availability
Observation mapping	120 scheduled; 10 new worlds	112 complete; 8 failures/interruption	Within-configuration disclosure effect
Action assays	45 cells + 360 slots	Open action + four conditions/model	Descriptive and failure-aware action transfer
Factorial intervention	160 slots; 10 worlds	120/120 sessions; 200/200 physics/replays	Fixed-evidence law and decision-rule replacement
Information separation	160 slots; 10 reused worlds	160/160 sessions; 80/80 physics/replays	Same-world artifact-only decision utility
Evaluator controls	16 unit versions	Provider-free; original stops retained	Rank validity versus action validity

4 RESULTS: FOUR OBSERVABLE CONVERSION QUESTIONS

4.1 F1: PRIOR-CONDITIONED SEARCH AND UNMET CORRECTION CRITERIA

DeepSeek produced records for all 135 scheduled sessions and completed 1,243/1,260 planned experiments; 121 sessions met operational eligibility. The denominator retains 26 discarded lifecycles, 13 resource-ledger rejections, and all right-censored cells. Every session submitted five checkpoints, providing 6,300 counterfactual predictions and 24,300 query–metric values.

Arm assignment was reflected in behavior. The first complete recipe differed between aligned and misspecified cells in 45/45 matched clusters, between opaque and aligned in 45/45, and between opaque and misspecified in 44/45. This is a manipulation check rather than a causal effect estimate because there are no repeated same-arm sessions. Search continued after the first proposal: 91.2% of completed experiments used a unique recipe, 84.4% of session optima appeared after the midpoint, and 32.6% appeared in the last completed experiment.

Correct-prior utility was task dependent. In entity-level partition, the aligned arm showed a +0.200 best-endpoint advantage over the misspecified arm in 5/5 worlds. In structural crystallization, a +0.141 first-experiment head start narrowed to +0.055 as the disadvantaged arm explored. In structural partition, aligned and misspecified descriptions both helped relative to opaque identifiers while differing little from one another. The supplied-model arms therefore occupied different search landscapes, without identifying a stochastic participant effect or imposing one endpoint ordering.

Prediction error nevertheless fell on average in every arm at every locus. Reductions for opaque, aligned, and misspecified cells were 0.111/0.097/0.097 at the entity locus, 0.090/0.033/0.065 at the parametric locus, and 0.219/0.228/0.221 at the structural locus. The stricter selective-correction contrasts were -0.214 for entity ($p = 0.990$), +0.033 for parametric ($p = 0.079$), and -0.224 for structural ($p = 1.000$); none passed. These criteria did not establish preferential correction of a wrong starting model (Fig. 2).

4.2 F2: STRUCTURE RECOVERY DEPENDS ON THE DISCLOSED MAPPING

Matched packets expose numerical–expression dissociation, but B2’s single-pair linear/power alias makes it an underidentifying surface. Historical B3 retained joint recovery of 5/30 for GPT and 0/30 for DeepSeek, including 13 DeepSeek schema failures. A later minimal-interface diagnostic completed 117/120 sessions but recovered structure in 0/80 unknown/wrong-prior cases. Optional




Figure 2: Prediction improves, but the evidence does not establish selective repair of the wrong model. a–c, DeepSeek and d–f, GPT effective pre/final prediction errors by prior and intervention locus; lines connect aggregate observed-point summaries with retained-checkpoint handling. Each model has 135 scheduled sessions in 45 world clusters. Confirmatory tests separately retain adverse failure bounds; all locus criteria are unmet in both models. Initial error limits improvement headroom. DeepSeek first-recipe differences are 45/45, 45/45 and 44/45 matched clusters; this retrospective check has no repeated same-arm baseline.

numerics changed success by +1.67 percentage points (95% interval [-8.33, 10.00]); GPT used the tool in 0/30 enabled sessions and DeepSeek in 20/30. The public mapping was incomplete, so these results do not establish failure on an identifiable task. Full historical denominators and controls remain in the appendix.

A separate disclosure intervention tests an actionable information condition. Ten new worlds span power exponents 1.35–2.25. Within each configuration, paired sessions receive the same observations, priors, minimal submission and optional calculator; only the post-evidence disclosure adds the complete observation mapping. GPT/medium and DeepSeek Flash/low each have 60 scheduled sessions. All 120 reached terminal states: 112 valid completions and eight DeepSeek failures or interruptions, retained in the denominator.

Among unknown or wrong priors, joint recovery rises from 6/20 to 20/20 for GPT and from 3/20 to 17/20 for DeepSeek (Fig. 3). Both gains are 70 percentage points; the pooled approximate world-bootstrap 95% interval is [52.5, 85]. Models and priors are averaged within each of ten shared worlds before resampling. Correct-prior retention is separate: 8/10 to 10/10 and 6/10 to 10/10, respectively.

Prediction error also decreases among available submissions, but action gains differ. GPT Top-1 rises from 5/20 to 12/20 and normalized regret falls from 0.07097 to 0.00748; DeepSeek Top-1 changes from 6/20 to 4/20 and regret from 0.15090 to 0.24263. The latter direction is descriptive, not a significant-worsening claim. Completing the mapping thus improves structural outputs in both tested configurations without guaranteeing action improvement. Disclosure length and paired observation noise are part of this setting; the new runtime and bounded power-family coverage do not retrospectively establish identifiability of historical packets. A selected one-retry sensitivity retains both directions: DeepSeek recovery 5/20 to 20/20 and Top-1 7/20 to 5/20; it does not replace the original primary.

4.3 F3: EXECUTABLE LAWS LOSE INFORMATION FROM PREDICTIONS

All 135 DeepSeek laws executed, but 84 lost information relative to final explicit predictions; mean law MAE was 0.237 and compression loss 0.069. Legal full-basis controls reproduced 135/135 prediction states at mean MAE 4.25×10^{-13} ; the schema can express these already-observed

270
271
272
273
274
275
276
277
278
279
280
281
282
283
284
285
286
287
288
289
290
291
292
293
294
295
296
297
298
299
300
301
302
303
304
305
306
307
308
309
310
311
312
313
314
315
316
317
318
319
320
321
322
323




Figure 3: Completing the observation mapping improves structure recovery; action gains differ. **a–b**, Joint family/exponent recovery and optimal action choice retain 20 scheduled unknown/wrong-prior sessions per configuration and disclosure condition. **c**, Prediction error uses the available submissions shown by *n*. **d**, Failed sessions receive worst normalized regret. All panels show the original attempts, with ten shared worlds; selected failure retries do not replace these outcomes. The two configurations are compared within system, not at equal compute.

prediction vectors. This same-domain fit does not test generalization. Blind incumbent replay completed 726 executions for 121 cells, with recommendations better/equivalent/worse in 1/119/1.


The matched GPT-5.6-sol surface retained 126 completed, 3 failed, and 6 right-censored cells. All locus gates again failed. GPT-5.6-sol versus DeepSeek-v4-flash law MAE was 0.1753 (129 laws) versus 0.2371 (135 laws), and compression loss was 0.0142 versus 0.0686; blind gain was -0.0001 (126 cells) versus -0.0010 (121 cells). Lower observed compression error therefore coexisted with near-zero blind gain (Fig. 4); this matched cross-configuration comparison is descriptive, not a causal law-quality intervention.

4.4 F4: SUBMITTED LAWS AND UNSEEN-PLAN DECISIONS CAN DISAGREE

The longitudinal DeepSeek cohort retained 45 scheduled cases and 42 terminal rankings, with 240/240 hidden plan evaluations and exact replays. Last-available laws selected the Top-1 plan in 0/45 cases, compared with 11/45 participant choices; participants followed their law in 12/42 available rankings. Failure-aware participant/law regret was 0.344/0.438. Law error had weak pooled association with regret, and task-specific associations reversed sign. An artifact is thus an incomplete behavioral proxy on this protocol; neither law following nor law quality was randomized, and the original cohort lacks a no-evidence action control.

A separate development successor retained four strategies over 45 strata per model (360 slots). Autonomy-minus-none regret was -0.0913 for DeepSeek (95% interval [-0.2124, 0.0388]) and +0.1102 for GPT ([-0.0533, 0.2794]). Both intervals cross zero and include donor and delivery failures; they do not isolate a pure acquisition effect. Completed-donor sensitivities, all failures and the descriptive law/action decomposition appear in the appendix (Fig. 8). Ranking-gate versus

Executable laws can lose information present in direct answers



Same law-evaluable cells at both endpoints: 135 DeepSeek / 129 GPT. Replay tests an observed plan.

Figure 4: **Lower executable-law error coexists with near-zero incumbent gain.** **a**, Hollow circles and filled squares show final-prediction and executable-law MAE on the same law-evaluable cells, grouped by locus and model; connecting lines show compression differences, not uncertainty. The matched denominators are 135 DeepSeek and 129 GPT laws. **b**, Better/equivalent/worse/unavailable incumbent-replay counts retain all 135 scheduled cells per model, with missing outcomes stated explicitly. These are descriptive configuration contrasts and incumbent replays, not causal artifact effects or unseen-action benefits.

decision-loss disagreement is separately an evaluator diagnostic, not a fourth internal participant fault (Table 6).

5 OBSERVED SUCCESSES: CONTROLLED ARTIFACT AND INFORMATION INTERVENTIONS

5.1 FIXED-EVIDENCE REPRESENTATION AND DECISION REPLACEMENT

The factorial assay fixes twelve public experiments per world and crosses model-generated (L) or public-ridge-fitted (F) quadratics with fresh-agent (A) or deterministic-maximizer (X) choice. Two tasks, five new worlds per task and two repeats per model yield forty source states nested within ten worlds. Sources never see the eight terminal candidates; recipients receive the evidence, candidates and designated law. All 120 sessions, 160 conditions and 200 physical executions with exact replay completed without failure or replacement.

F-X minus L-X regret was -0.00538 (95% world-bootstrap interval [-0.01630, 0.00061]); the upper endpoint did not fall below the prespecified -0.01 material-benefit threshold. Task means were -0.01087 in electrochemistry and +0.00010 in crystallization, with the substantial negative effect concentrated in one world. Figure 9 in the appendix shows all world effects and five registered contrasts. The fixed utility scale is one; models and repeats are averaged within world, and five worlds per task limit the bootstrap approximation.

Agent/maximizer choices agreed in 39/40 model-law and 40/40 fitted-law pairs. The latter's zero-width bootstrap interval describes observed equality, not population equivalence. Fitted-maximizer regret was 0.00425, versus 0.00354 for nearest public evidence and 0.11109 for uniform-random expected choice. There was no observed advantage over retrieval. This bounds a general law-use failure account, but the co-delivered raw evidence prevents an artifact-only interpretation; the causes of historical cross-protocol disagreement remain unidentified.

5.2 INFORMATION SEPARATION REVEALS INDEPENDENT ARTIFACT VALUE

We reused all ten factorial worlds and forty sealed sources, fixed eight new candidates per world, and assigned fresh recipients task information alone, raw evidence, model law alone (L) or fitted




Figure 5: Artifacts support fresh decisions, while retrieval remains a strong baseline. **a**, Dots show world means over two models and two repeats per information condition; diamonds show task means. All 160 selections are available. Nearest evidence has zero regret in all ten worlds. **b**, Colored points show paired world effects; black diamonds and lines show means and prespecified intervals (95% primary; 99% for five secondary comparisons). The dotted primary reference is -0.01. L/F denote model-generated/fitted laws supplied without raw observations. These are new candidate plans in the ten reused worlds, not new physical mechanisms or additional independent replication worlds.

law alone (F). Only the raw condition received observations; law conditions received six original coefficients without source dialogue or provenance labels. All 160 sessions and 80 hidden evaluations with exact replay completed, with no new recipient measurements.

The primary L-minus-none regret contrast was -0.13723 (95% interval [-0.15584, -0.12257]), meeting the prespecified material-benefit criterion (Fig. 5). Nine world means improved and one was zero; task means were -0.25937 for electrochemistry and -0.01509 for crystallization. These remain ten reused worlds, not ten additional independent replications. Mean regrets for none/raw/L/F were 0.14727/0.01124/0.01004/0.01459. Raw and F also improved on none under adjusted intervals. L minus raw was -0.00120 (99% interval [-0.02353, 0.00951]); this establishes neither superiority nor equivalence. Nearest public evidence selected the measured optimum in all ten worlds, so these data do not establish an advantage over retrieval.

Law-only deployment used 392,101 input tokens versus 541,225 for raw evidence, a descriptive 27.6% reduction; source acquisition and generation costs had already been incurred. Model-law recipients followed the deterministic maximizer in 40/40 states and fitted-law recipients in 39/40. The result supports same-world context portability of the delivered artifact, without identifying internal law use, transfer across changed mechanisms or experimental savings.

6 RELATED WORK

Self-driving laboratories and chemistry agents emphasize closed-loop execution, tool use, or endpoint optimization (Felton et al., 2021; Häse et al., 2021; Boiko et al., 2023; Bran et al., 2024). Virtual laboratories and process-control environments provide scalable interaction and safety (Beeler et al., 2024; Bloor et al., 2024; Malik et al., 2026; Chen et al., 2025). Scientific-discovery benchmarks increasingly test iterative experimentation, causal inference, and transferable knowledge (Jansen et al., 2024; Gandhi et al., 2025; Duan et al., 2025; Yang et al., 2026; Batzoglu, 2026). The published ChemWorld platform contributes programmable chemical worlds, transaction semantics

and replay (Qiu et al., 2026); this work contributes the intervention and measurement programme for scientific-agent epistemics.

Decision-focused learning distinguishes prediction error from downstream decision loss (Elmachtoub & Grigas, 2022; Wilder et al., 2019). Our additional setting is an agent that acquires evidence under operational constraints and submits a reusable knowledge artifact. We measure both the fidelity of that artifact to the agent’s own predictions and its agreement with actual choices, then deliver it separately from source observations to test recipient decision utility. These are distinct empirical tests; we do not introduce a decision-focused training algorithm.

Model-discovery systems increasingly combine language models with Bayesian design, symbolic fitting or probabilistic programme search (Murphy, 2026; Wahl et al., 2026; Zheng et al., 2026). We instead ask whether evidence changes the right representation and whether that representation is usable for action. This follows the broader distinction between outcome and process validity (Ríos-García et al., 2026) and the warning that predictive success can coexist with underspecified or shortcut solutions (D’Amour et al., 2022; Geirhos et al., 2020). Causal mediation analysis would require identified interventions on intermediate representations (Imai et al., 2010); we do not infer such mediation from associated law and action readouts. The central contribution is therefore a tiered diagnostic design: controlled initial-model manipulations, conditional packet responses, failure-aware strategy estimates, descriptive law/action decomposition, a fixed-evidence factorial intervention, and information-separated artifact deployment with separate evaluator qualification.

7 DISCUSSION AND LIMITATIONS

These studies distinguish information availability, artifact fidelity and decision value. Completing observation mapping improves structure in both configurations, while action gains differ; structural accuracy cannot substitute for measuring the resulting decision. Historical prediction gains did not establish selective repair, and submitted laws often differed from participant choices. The fixed-evidence intervention found high agent/maximizer agreement without a supported material fitted-law benefit. Information separation then established that a compact model law can help a fresh recipient on new plans in the same worlds. Raw evidence also helped, and nearest retrieval achieved zero regret; this is a bounded knowledge-utility result, without demonstrated superiority or a new repair algorithm.

The information comparison evaluates deployment strategies, including context quantity and computation. Matching recipient tools does not remove the original tool-free-model versus numerical-ridge difference at source. Five worlds per task support approximate intervals, and reuse does not add independent replication worlds. Each task shares a candidate grid across its five worlds, and the measured optimum has the same coordinates within task. Effects differ by task. The protocols also differ in information, function class, decision dimension and interface; their contrast does not isolate a single cause or internal psychological mediation. Private confirmation, new-physical-condition transfer and independent-backend replication remain untested.

The disclosure effect covers ten shared worlds in one power family, with paired noise and additional prompt length. DeepSeek/low and GPT/medium are study-specific configurations, not equal-compute competitors; calendar and concurrency amendments are disclosed. Historical controls retain additional limits. B2 is underidentifying; historical B3 retains 13 DeepSeek schema failures. The minimal-interface diagnostic retains three failures and a tool effect interval crossing zero; absent GPT tool uptake limits any arithmetic interpretation. Its privileged reference qualification does not establish public-only identifiability. Low reasoning is not thinking-off. Donor-eligible analyses are sensitivity-only, and fixed order and earlier donors confound configuration with time. The earlier four-condition strategy study retains delivery failures and cannot isolate pure acquisition. Decision loss, near-optimality, availability and source/deployment costs should therefore remain separate; best-minus-first gains alone do not establish feedback learning or experimental savings.

AI USE STATEMENT

Generative AI tools assisted software implementation, experiment orchestration, data analysis, visualization code, literature discovery, and drafting and editing of the manuscript. The authors reviewed the study designs and frozen decision rules, inspected retained failures, verified reported denominators and claims against machine-readable records, and tested AI-assisted code before use. The authors take responsibility for the final content, including text, claims, code, and artifacts produced with generative-AI assistance.

REPRODUCIBILITY STATEMENT

The main paper defines the experimental units, interventions, estimands, failure handling, and stop boundaries. The appendix provides the execution and evaluation details needed to reconstruct each evidence layer. Sanitized scheduled-cell records, evaluator-derived metrics, replay and source digests, retained failure classifications, protocol projections, prompts, schemas, and figure source tables are included in the anonymous supplementary package.

ETHICS STATEMENT

The experiments use simulated chemical environments and no human participants, personal data, or live laboratory actuation. The principal risks are over-interpreting simulated results and treating an evaluator surrogate as the scientific objective; the claim boundaries and retained negative qualification results are reported to mitigate these risks.

REFERENCES

- Serafim Batzoglou. ReplaySCM: A benchmark for executable causal mechanism induction from interventions, 2026. Preprint.
- Chris Beeler, Sriram Ganapathi Subramanian, Kyle Sprague, et al. ChemGymRL: A customizable interactive framework for reinforcement learning for digital chemistry. *Digital Discovery*, 3: 742–758, 2024. doi: 10.1039/D3DD00183K.
- Maximilian Bloor, José Torracca, Ilya Orson Sandoval, et al. PC-Gym: Benchmark environments for process control problems, 2024. Preprint, arXiv:2410.22093.
- Daniil A. Boiko, Robert MacKnight, Ben Kline, and Gabe Gomes. Autonomous chemical research with large language models. *Nature*, 624:570–578, 2023. doi: 10.1038/s41586-023-06792-0.
- Andres M. Bran, Sam Cox, Oliver Schilter, Carlo Baldassari, Andrew D. White, and Philippe Schwaller. Augmenting large language models with chemistry tools. *Nature Machine Intelligence*, 6:525–535, 2024. doi: 10.1038/s42256-024-00832-8.
- Yimeng Chen, Piotr Piękos, Mateusz Ostaszewski, Firas Laakom, and Jürgen Schmidhuber. PhysGym: Benchmarking llms in interactive physics discovery with controlled priors. In *Advances in Neural Information Processing Systems*, 2025. URL <https://openreview.net/forum?id=w8uII2qAmd>.
- Alexander D’Amour, Katherine Heller, Dan Moldovan, Ben Adlam, Babak Alipanahi, Alex Beutel, Christina Chen, Jonathan Deaton, Jacob Eisenstein, Matthew D. Hoffman, Farhad Hormozdiari, Neil Houlsby, Shaobo Hou, Ghassen Jerfel, Alan Karthikesalingam, Mario Lucic, Yian Ma, Cory McLean, Diana Mincu, Akinori Mitani, Andrea Montanari, Zachary Nado, Vivek Natarajan, Christopher Nielson, Thomas F. Osborne, Rajiv Raman, Kim Ramasamy, Rory Sayres, Jessica Schrouff, Martin Seneviratne, Shannon Sequeira, Harini Suresh, Victor Veitch, Max Vladymyrov, Xuezhi Wang, Kellie Webster, Steve Yadlowsky, Taedong Yun, Xiaohua Zhai, and D. Sculley. Underspecification presents challenges for credibility in modern machine learning. *Journal of Machine Learning Research*, 23(226):1–61, 2022. URL <https://www.jmlr.org/papers/v23/20-1335.html>.

- Kourosh Darvish, Marta Skreta, Yuchi Zhao, et al. ORGANA: A robotic assistant for automated chemistry experimentation and characterization. *Matter*, 8(2), 2025. doi: 10.1016/j.matt.2024.10.015.
- Haonan Duan, Stephen Zhewen Lu, Caitlin Fiona Harrigan, Nishkrit Desai, Jiarui Lu, Michał Koziarski, Leonardo Cotta, and Chris J. Maddison. Measuring scientific capabilities of language models with a systems biology dry lab, 2025. Preprint, arXiv:2507.02083.
- Adam N. Elmachtoub and Paul Grigas. Smart Predict, then Optimize. *Management Science*, 68(1): 9–26, 2022. doi: 10.1287/mnsc.2020.3922.
- Kobi C. Felton, Jan G. Rittig, and Alexei A. Lapkin. Summit: Benchmarking machine learning methods for reaction optimisation. *Chemistry–Methods*, 1:116–122, 2021. doi: 10.1002/cmtd.202000051.
- Kanishk Gandhi, Michael Y. Li, Lyle Goodyear, Agam Bhatia, Louise Li, Aditi Bhaskar, Mohammed Zaman, and Noah D. Goodman. BoxingGym: Benchmarking progress in automated experimental design and model discovery, 2025. Preprint, arXiv:2501.01540.
- Robert Geirhos, Jörn-Henrik Jacobsen, Claudio Michaelis, Richard Zemel, Wieland Brendel, Matthias Bethge, and Felix A. Wichmann. Shortcut learning in deep neural networks. *Nature Machine Intelligence*, 2:665–673, 2020. doi: 10.1038/s42256-020-00257-z.
- Florian Häse, Matteo Aldeghi, Riley J. Hickman, Loïc M. Roch, Melodie Christensen, Elena Liles, Jason E. Hein, and Alán Aspuru-Guzik. Olympus: a benchmarking framework for noisy optimization and experiment planning. *Machine Learning: Science and Technology*, 2:035021, 2021. doi: 10.1088/2632-2153/abedc8.
- Kosuke Imai, Luke Keele, and Dustin Tingley. A general approach to causal mediation analysis. *Psychological Methods*, 15(4):309–334, 2010. doi: 10.1037/a0020761.
- Peter Jansen, Marc-Alexandre Côté, Tushar Khot, Erin Bransom, Bhavana Dalvi Mishra, Bodhisattwa Prasad Majumder, Oyvind Tafjord, and Peter Clark. DiscoveryWorld: A virtual environment for developing and evaluating automated scientific discovery agents. In *Advances in Neural Information Processing Systems*, volume 37, 2024. doi: 10.52202/079017-0324.
- Shreshth A. Malik, Tiarnan Doherty, Panagiotis Tigas, Muhammed Razzak, Stephen J. Roberts, Aron Walsh, and Yarin Gal. MADE: Benchmark environments for closed-loop materials discovery. In *International Conference on Machine Learning*, 2026.
- Kevin Murphy. Model discovery agent: Llm-assisted bayesian experiment design for data-efficient discovery of mechanistic world models, 2026. Preprint, arXiv:2608.09696.
- Jiangjie Qiu, Yijun Li, and Xiaonan Wang. ChemWorld: Programmable chemical worlds for controlled and replayable agent experimentation, 2026. Preprint, arXiv:2608.10792.
- Martíño Ríos-García, Nawaf Alampara, Chandan Gupta, Indrajeet Mandal, Sajid Mannan, Ali Asghar Aghajani, N. M. Anoop Krishnan, and Kevin Maik Jablonka. Ai scientists produce results without reasoning scientifically, 2026. Preprint, arXiv:2604.18805.
- Tao Song, Man Luo, Xiaolong Zhang, et al. A multiagent-driven robotic AI chemist enabling autonomous chemical research on demand. *Journal of the American Chemical Society*, 147(15): 12534–12545, 2025. doi: 10.1021/jacs.4c17738.
- Nathan J. Szymanski et al. An autonomous laboratory for the accelerated synthesis of inorganic materials. *Nature*, 624:86–91, 2023. doi: 10.1038/s41586-023-06734-w.
- Aikaterini Vriza, Michael H. Prince, Tao Zhou, Henry Chan, and Mathew J. Cherukara. Operating advanced scientific instruments with AI agents that learn on the job. *npj Computational Materials*, 12:160, 2026. doi: 10.1038/s41524-026-02005-0.
- Stefan Wahl, Raphaela Schenk, Ali Farnoud, Jakob H. Macke, and Daniel Gedon. A probabilistic framework for llm-based model discovery, 2026. Preprint, arXiv:2602.18266.

594 Bryan Wilder, Bistra Dilkina, and Milind Tambe. Melding the data-decisions pipeline: Decision-
595 focused learning for combinatorial optimization. In *Proceedings of the AAAI Conference on*
596 *Artificial Intelligence*, volume 33, pp. 1658–1665, 2019. doi: 10.1609/aaai.v33i01.33011658.
597

598 Junlin Yang, Dylan Zhang, Xiangchen Song, et al. CausaLab: A scalable environment for interactive
599 causal discovery toward ai scientists, 2026. Preprint.

600 Tianshi Zheng, Kelvin Kiu-Wai Tam, Newt Hue-Nam K. Nguyen, et al. NewtonBench: Benchmarking
601 generalizable scientific law discovery in llm agents. In *International Conference on Learning*
602 *Representations*, 2026.
603
604
605
606
607
608
609
610
611
612
613
614
615
616
617
618
619
620
621
622
623
624
625
626
627
628
629
630
631
632
633
634
635
636
637
638
639
640
641
642
643
644
645
646
647

A APPENDIX

A.1 ENVIRONMENT SEMANTICS AND INTERVENTION CONSTRUCTION

An executable world is a stateful transition system with typed operations, measurement functions, finite resources, terminal assays, and an evaluator-only truth surface. A task publishes the operation and instrument contracts while withholding the latent state and registered relations. A scenario binds a world, public task description, resource card, and initial model. Initial-model arms alter one declared information locus without changing the transition process, measurement process, candidate actions, or evaluator truth.

The entity locus changes material identity or ontology. The parametric locus changes a quantitative response model while preserving the declared entities. The structural locus changes the functional form used to describe a mechanism. Aligned and misspecified models are matched in explicitness so that a contrast is not merely information versus no information. The opaque arm exposes the same laboratory interface but withholds the intervened model component. Exploratory, validation, prospective, and private world selections are disjoint by construction.

A.1.1 TASK-LOCUS ROSTER AND WORLD SPLITS

The nine prospective task-locus combinations are fixed before participant execution. The entity intervention swaps matched nominal descriptor rows while leaving physical identity fixed. Parametric and structural interventions replace only the declared response component; all other task text, tools, resources, queries, and world state remain matched.

Table 2: Nine prospective task-locus combinations. “Registered target” names the component manipulated in the initial model, not information revealed by the evaluator.

Locus	Task	Registered target	Held-out diagnostic
Entity	Electrochemical conversion	solvent identity descriptors	conversion/selectivity response
Entity	Reaction-crystallization	solvent identity descriptors	reaction and crystal outcomes
Entity	Reaction-distillation	solvent identity descriptors	reaction and separation outcomes
Entity	Partition discovery	extractant identity descriptors	phase-partition outcomes
Entity	Reaction safety	catalyst identity descriptors	yield, kinetics, and risk
Parametric	Electrochemical conversion	quantitative potential-response direction	peak-and-collapse profile
Parametric	Reaction safety	quantitative catalyst/thermal response	productive versus hazardous regime
Structural	Partition discovery	linear versus power-form phase relation	registered exponent 1.75
Structural	Reaction-crystallization	pathway/topology and accumulation form	multistage outcome surface

Development and qualification use seeds 0–4 and never enter the prospective denominator. Public worlds are five task-specific SHA-256-derived identities per task; their values are fixed in the released configuration. The private split is bound by a commitment but was not unsealed or executed. Matched-evidence and action studies use separately frozen public surfaces. Sessions sharing a task-world but differing in prior arm are repeated measurements within one inference cluster.

A.1.2 METRICS AND REGISTERED CORRECTION GATES

For prediction or executable law f , truth y , registered query-metric pairs $j = 1, \dots, J$, and fixed positive metric scale s_j (one for bounded metrics unless overridden), normalized error is

$$E(f) = \frac{1}{J} \sum_{j=1}^J \min\left(\frac{|f_j - y_j|}{s_j}, 1\right).$$

For candidate scores q_i , selected candidate a , best $q_{\max}$, and worst $q_{\min}$, normalized regret is $(q_{\max} - q_a)/(q_{\max} - q_{\min})$, with zero assigned when the packet range is numerically zero. A missing or invalid terminal ranking receives failure-aware regret one and Top-1 zero. Near-optimal selection means raw regret at most 0.01; pairwise ordering excludes truth pairs separated by less than 0.01.

Let $I_a = E_{a,0} - E_{a,K}$ and $C = I_{\text{mis}} - I_{\text{aligned}}$. The entity gate is a three-component intersection-union test at one-sided $\alpha = 0.05$: the lower bound for C must exceed zero, misspecified improvement

must exceed zero, and aligned improvement must exceed the non-inferiority margin -0.05 . Parametric and structural gates use equal-task fixed-effect estimates of C , require the one-sided lower bound above zero and positive direction in every included task. For an incomplete cell, confirmatory bounds use the adverse interval $[-1, +1]$ for arm improvement; the corresponding contrast can therefore span $[-2, +2]$. Last-observation-carried-forward and zero improvement are observed-point sensitivities only and do not generate the confirmatory p values.

A.2 PERSISTENT EXECUTION AND CHECKPOINTS

Each participant cell is one persistent agent process. At every turn, the agent may call only the public laboratory tools. The host validates schemas and preconditions, applies resource accounting, commits successful transactions, and returns public observations. It does not choose or repair an experiment on the participant’s behalf. Entity, parametric, and structural campaigns permit eight, ten, and twelve complete experiments, respectively. Checkpoints occur at experiments 0/2/4/6/8, 0/2/4/7/10, and 0/3/6/9/12.

Every checkpoint records model-reliability judgements, query predictions, uncertainty, public evidence references, a typed executable law, and the intended next experiment. Checkpoint repair is part of the evaluated interface burden. All eventually yielded valid records, but rejected attempts remain in the process log and resource accounting.

A.2.1 SYSTEM, PROMPT, TOOL, AND BUDGET CONTRACT

`DeepSeek-v4-flash` with `reasoning_effort=high` supplies the primary DeepSeek C2, matched-evidence, B3-successor, and action surfaces; `low` supplies the completed B2 reasoning-budget ablation and the later observation-mapping disclosure block. `GPT-5.6-sol` with `reasoning_effort=medium` supplies the matched GPT replications. Both run through the Codex session harness; “Codex” therefore denotes execution infrastructure, not a third participant model. Provider aliases, run dates, receipt counts, and usage fields remain in the retained internal records rather than the anonymous supplement. The supplement retains sanitized model/configuration labels, scheduled completion denominators, and derived resource summaries. Temperature is not exposed or set by the harness. Request timeout is 1,200 s; finalization timeout is 600 s; progress is emitted every 30 s.

Each C2 cell receives one task/scenario prompt, the same public operation schemas, the arm-specific initial model, and no evaluator truth. A persistent transcript spans the complete campaign. The fixed 8/10/12-experiment budgets allow at most 96/110/168 operation attempts, 480k/600k/720k uncached input tokens, 128k/160k/192k output tokens, and 7,200/7,200/9,000 s wall time, respectively. Each cell uses at most two provider calls within the same session; infrastructure-only resume is capped once, while any persisted scientific trajectory forbids replacement. Matched-evidence assays use exactly two same-thread turns: a pre-evidence submission and a post-packet submission. The supplement includes the exact static system prompts, selected prompt-generating function sources, public protocol projections, representative strict response schemas, and the executable-law contract. The package also contains all 45 sanitized B2 public summaries, the retrospective expression-coding function, three-configuration all-arm counts, and the participant-visible B2 identifiability audit.

A.2.2 MODEL COVERAGE AND THE UNIT OF INFERENCE

The programme developed from a DeepSeek cohort to matched GPT successor studies and later two-system interventions. The included systems are not a representative sample of model families. A configuration includes the model, reasoning setting, prompts, tools and execution policy. A second model checks recurrence of a within-system effect; shared worlds remain shared inference units, and the common harness limits external validity.

Interventions compare information, artifact source or decision rule within a configuration. Across-model comparisons are descriptive: equal laboratory budgets do not equate hidden reasoning compute, and effort labels are not a shared scale. Historical settings and prespecified pooled weights are unchanged. Pooled effects describe these systems only; different significance decisions do not establish a model-by-intervention interaction or equivalence. Failed model integrations provide engineering evidence, not additional scientific participants.

Table 3: **Model coverage follows the scientific question.** Counts are scheduled sessions; failures remain in each protocol’s denominator.

Scientific question	Executed model coverage	Inferential role
Search, correction and law fidelity	C2: DeepSeek and GPT, 135 sessions each	Matched surface; original gates retained
Evidence and interface diagnostics	Both families; extra DeepSeek-low B2 ablation	Separate protocol-specific controls
Longitudinal law/action agreement	Original 45-session DeepSeek cohort only	Single-system descriptive evidence
Law replacement and isolated delivery	Both families in every M1/M3 condition	Fixed system mixture in ten worlds
Observation-mapping disclosure	GPT/medium and DeepSeek/low, 60 sessions each	Within-system effects in ten shared worlds

A.2.3 EXECUTABLE-LAW SCHEMA AND ACTIONPLAN CONSTRUCTION

A law contains one metric equation for every required metric, an intercept, identity or logistic link, output bounds, evidence references, confidence, applicability, and limitations. Each metric permits at most 64 terms drawn from linear, quadratic, cubic, pairwise interaction, categorical indicator, and category-conditional linear/quadratic/cubic bases. Terms may reference only the task’s registered feature IDs; evidence references may name only participant-visible assays. The production parser and executor score the law directly, so prose that cannot be parsed does not silently receive a law score.

Terminal candidates are eight complete, compiler-valid ActionPlans selected from a registered 16-query public-feature roster without reading outcomes. Each packet declares every ordered operation and parameter, fresh initial state, measurement positions, terminal assay, omitted optional operation, and objective. The other eight rows form the disjoint prediction/oracle-fit roster. Public-plan, truth-plan, and replay hashes must agree exactly. Candidate outcomes and evaluator ranks remain hidden until scoring.

A.3 EVALUATOR-OWNED PREDICTION AND LAW ASSAYS

Entity checkpoints contain four prespecified counterfactual queries; parametric and structural checkpoints contain sixteen. The evaluator executes each unique task–world query independently, never returns its truth to the participant, and shares truth across prior arms and checkpoints. The primary readout is mean normalized absolute error across registered query–metric pairs. The executed cohort contains 420/420 truth executions, 1,620 query–metric truth values, 675/675 checkpoints, 6,300 predictions, and 24,300 scored values.

Final typed laws are executed on the same registered coordinates. Schema validity, coordinate coverage, truth-normalized error, change from pre-evidence and effective-final checkpoints, and consistency with explicit final predictions are retained separately. This assay establishes executability and in-domain compression only. Reusability would require new coordinates or transfer to a context-reset process.

The matched GPT-5.6-sol evaluator uses the identical 45-cluster, nine-task, three-arm C2 surface. DeepSeek-v4-flash and GPT-5.6-sol have 121/135 and 126/135 completed cells. Their exact evaluator denominators are 675/675 versus 669/675 scored checkpoints, 135/135 versus 129/135 executable laws, and 121/135 versus 126/135 cells with evaluable blind gain. Overall mean prediction improvement is 0.1198/0.1329, final prediction error is 0.1685/0.1614, law MAE is 0.2371/0.1753, law compression loss is 0.0686/0.0142, and blind gain is -0.0010/-0.0001. GPT-minus-DeepSeek paired descriptive differences are +0.0131 for prediction improvement, -0.0071 for final error, -0.0648 for law MAE, and -0.0579 for compression loss. Both models fail all three registered selective-correction gates.

A.4 TYPED-LAW CAPACITY AND DISTILLATION CONTROLS

To distinguish participant compression failure from an underpowered output schema, we refitted each of the 135 complete final prediction vectors with legal identity-link laws over the registered feature coordinates and allowed basis functions. A full-basis fit used up to 64 terms per metric. A term-matched fit used exactly the participant’s submitted term budget for each metric. A leave-one-query-out fit withheld each registered query in turn and predicted it from all remaining queries. Every fitted law was parsed and executed by the production law evaluator; independently computed design-matrix predictions had to agree with executor output within 10^{-10} .

The participant laws differed from their own final predictions by mean MAE 0.1539. Full-basis fits reproduced 135/135 final prediction states with mean MAE 4.25×10^{-13} ; term-matched fits reduced the error to 0.0114 and leave-one-query-out fits reached 0.0788. The control uses the same registered coordinate domain and therefore tests representation capacity and distillation, not global mechanistic identification or transfer.

A.5 FAILURE-AWARE INFERENCE

The unit of prospective inference is the independently selected task–world cluster, not a checkpoint, query, experiment, or replay. The three prior arms within a cluster are matched. Locus-specific selective-correction contrasts are computed first and are not pooled across different intervention semantics. Failed scientific cells remain assigned. Confirmatory correction bounds assign an incomplete post-outcome the adverse improvement range $[-1, +1]$; last-observation-carried-forward and zero improvement are observed-point sensitivities only. Infrastructure recovery is allowed only when no scientific trajectory has persisted and only under the fixed attempt limit.

The matched-evidence B2 contrast has five independent worlds. Its interval and exact one-sided sign-flip result are descriptive because the block was designed to localize a transition, not to establish a population-level arm effect. Recipe divergence is likewise a manipulation check: without repeated same-arm sessions, provider stochasticity cannot be separated from intervention sensitivity at exact-recipe resolution.

DeepSeek and GPT-5.6-sol medium used identical matched-evidence worlds, arms, evidence packets, queries, and scoring rules. Each completed 15/15 A-P and 15/15 A-S B2 formal sessions with no failed cells; canaries were excluded. A-P misspecified-minus-aligned update contrasts were 0.0309 and 0.0602, with 3/5 and 5/5 positive worlds. A-S B2 contrasts were 0.0645 and 0.0915, with 3/5 and 4/5 positive worlds; misspecified exact 1.75-law expression was 0/5 in both models. Configurations are reported separately, and no model-superiority test is performed.

The B2 exact-law counts are retrospective keyword coding of public `model_summary` and `evidence_assessment` fields; no typed family or exponent field was preregistered. A later participant-visible audit found that evidence and scoring fixed one nominal solvent/extractant pair, did not expose the base partition coefficient, and admitted an exact alias between the 1.75-power law and a free-coefficient linear law (coefficient multiplier 3.13588). A constant endpoint baseline obtained mean scoring MAE 0.00649, while the aligned DeepSeek-high exact-law positive control was only 1/5 and failed its readout criterion. B2 therefore supports conditional low post-packet error and lack of stable exact-law expression on an underidentifying surface, not a participant-level structural-identification failure. The separate B3 control below supplies a typed structural readout; its privileged fitter does not establish participant-visible information sufficiency.

The reasoning-budget ablation kept the DeepSeek model, Codex harness, prompts, schemas, public packets, worlds, and evaluator fixed while changing only `reasoning_effort` from high to low. It is not provider-level thinking-off, which would require a different direct-controller harness. The B2 interface canary passed 3/3 and the formal block completed 15/15 sessions, 30/30 turns, 15/15 same-thread continuations, and 360/360 pre plus 360/360 post scoring terms, with no failed session or infrastructure predecessor. Mean post errors were 0.0067/0.0069/0.0069 for opaque/aligned/misspecified. The registered contrast was -0.0405, with 2/5 positive worlds, exact one-sided $p = 0.8125$, and descriptive interval $[-0.1559, 0.0749]$. Misspecified exact 1.75-law expression remained 0/5; all five misspecified post errors were at most 0.02. Aligned exact-law expression was 1/5 for DeepSeek high, 0/5 for GPT medium, and 2/5 for DeepSeek low, reinforcing that the free-text expression readout itself lacks a strong positive control.

Provider-reported DeepSeek reasoning output was 506,637 tokens at high effort and 400,639 at low effort on the matched B2 block, a 20.9% reduction. These resource fields are descriptive within the DeepSeek provider and are not compared with GPT accounting. The separate low-effort A-P canary was platform-defective: progress events existed, but terminal cell receipts and a canary summary did not. It is retained as an unscored partial; A-P low formal execution remained 0/15 and supplies no effect estimate.

A.6 REFERENCE-FITTER-IDENTIFIABLE LAW AND ACTION CONTROL

The independent B3 control used five frozen public structural-partition worlds, the same opaque, aligned, and misspecified prior arms, and two independent fresh GPT-5.6-sol medium sessions per arm and world. Provider-free development qualification selected eight evidence rows spanning four nominal pairs and a disjoint eight-query scoring/action roster. Public truth was executed before any participant call. All five worlds entered family, exponent, prediction, Top-1, rank, and regret denominators; action gain was scored only in the three worlds whose best candidate exceeded the visible evidence incumbent by at least 0.02. The evidence and scoring rosters, five worlds, three arms, two replicates, 0.10 exponent tolerance, and 0.02 opportunity threshold were frozen before participant execution.

A three-session canary tested only two-turn schema validity, packet identity, same-thread continuity, typed law fields, the exact scoring denominator, and selection of an unexecuted candidate. It passed 3/3 and did not enter the formal scientific denominator. The full 30-cell block then completed 30/30 with no failed cells or replacement. All sessions preserved same-thread continuity and returned 60/60 completed formal turn receipts. The canary and formal blocks together used 33 provider session attempts, 66/66 completed turn receipts, zero retries, zero tool events, and zero participant physical experiments. One formal post turn recorded a transient provider error event before returning a completed receipt; the event remains in the resource ledger and the cell remains completed.

Joint structural recovery required both the registered power family and exponent error at most 0.10. Counts for opaque, aligned, and misspecified arms were 0/10, 5/10, and 0/10. The corresponding mean post-evidence errors were 0.0367, 0.0215, and 0.0378. The aligned world-mean exponent error was lower than each comparator in all five worlds. Top-1 counts were 0/10, 2/10, and 0/10; both Top-1 choices came from one action-ineligible world. None of the 18 action-eligible cells reached gain 0.02, including neither of the two action-eligible cells with joint structural recovery. A matched DeepSeek B3 successor then started from its first cell on the same scientific surface and retained every scheduled outcome. It completed 17/30 cells; 13/30 were participant-schema failures. Failure-aware joint recovery was 0/30, Top-1 was 0/30, and mean regret was 0.9579; completed-cell post MAE was 0.0928. The corresponding GPT-5.6-sol values were 5/30, 2/30, 0.7594, and 0.0320. Useful-gain success was 0/13 versus 0/18 among completed evaluable opportunities and 0/18 for both models on the fixed scheduled-opportunity denominator, with failures and unavailable rows counted as zero. The matched comparison is descriptive because model configuration was not randomized and schema-failure patterns differed.

A.7 MINIMAL-INTERFACE AND OPTIONAL-NUMERICS DIAGNOSTIC

The final diagnostic reuses five scored partition worlds and their sealed public packets, three initial descriptions, two model configurations and two fresh-session repeats. Within each combination, post-evidence numerics is either unavailable or optional, giving 120 sessions and 240 pre/post turns. The pre stage has no tools; post resumes the same thread. Family, exponent and eight query predictions are declared once; post also selects a query. The public mapping to a typed law does not fit parameters or inspect hidden truth. No new independent worlds, participant physics, truth execution or replay were produced.

The numerical interface accepts only literal array arithmetic and a fixed whitelist including logarithms and least squares. It cannot read files, use the network or invoke the simulator. At most eight attempted calls are permitted per session; rejected expressions also consume that allowance. Budgets were fixed after a separate twelve-session development block: 600 seconds per turn, 1,200 per session and 18 hours overall, with request and stream retries disabled. GPT used medium reasoning and DeepSeek high; these are not equal-compute settings.




Figure 6: **Numerical fit and structural recovery require different tests.** **a**, All 15 DeepSeek-high B2 cells before and after the same packet, shown on a log scale; gray/teal/orange denote opaque/aligned/misindexed arms. All post-packet errors are below 0.02, but the one-pair surface underidentifies family; low error is not structural recovery. **b**, B3 completion, joint family–exponent recovery, Top-1, and useful gain retain all scheduled opportunities. Counts distinguish 30 sessions from 18 gain opportunities per model, including 13 DeepSeek schema failures. Configuration contrasts are descriptive.


Table 4: **Final diagnostic outcomes retain every scheduled session.** Each model/tool group has 30 sessions and 18 prespecified useful-gain opportunities. Prediction MAE uses the complete paired denominator in column two.

Model / numerics	Valid	Pre / post MAE	Joint	Useful gain
GPT / off	29/30	0.2963 / 0.0367	6/30	0/18
GPT / on	30/30	0.2891 / 0.0435	7/30	2/18
DeepSeek / off	29/30	0.2946 / 0.0695	2/30	0/18
DeepSeek / on	29/30	0.2899 / 0.0718	2/30	0/18

Joint success requires the power family and an exponent within 0.10 of 1.75. A failed session scores zero joint success and normalized regret one; predictive errors use explicit available-pair denominators. The sole primary contrast is tool-on minus tool-off joint success, averaging priors and repeats within model and the two models equally within world. The five world means are resampled 20,000 times with seed 90770 for a percentile 95% interval. The interval is a small-sample approximation, not evidence that 120 sessions are independent worlds. Aligned priors already supply the correct family and exponent: their success measures retention.

The terminal block contains 117 valid sessions and three retained failures: one DeepSeek tool-enabled post-submission validation failure, one GPT transport interruption and one DeepSeek provider output-limit failure. All 240 turns were attempted, without replacement. The separate development block had ten valid sessions and two provider failures; it is excluded from formal denominators. The five paired world effects were +0.16667, -0.16667, 0, 0 and +0.08333. Opaque and misspecified strata contain zero joint recoveries in 80 scheduled sessions; the 17 aligned successes measure retention. Top-1 choices were 0/30, 1/30, 0/30 and 0/30 in table order. These action readouts are restricted query choices, not a replication of the longitudinal assay.

Total recorded wall time was 33,619 seconds (9.34 hours). Reported usage was 8,251,168 input and 5,152,198 output tokens; two failed turns lack usage, so these are known lower bounds. GPT made no numerical calls in its 30 enabled sessions. DeepSeek attempted 63 calls in 20/30 enabled sessions, including nine rejected expressions; tool computation totaled 0.027 seconds. Availability includes



10 sessions per prior/model/tool; 117/120 valid. Failures retained.

*Aligned supplies the law: retention. Five reused worlds; approximate interval.

Figure 7: **Valid submissions and optional numerics do not ensure structural recovery.** **a**, Joint family/exponent success by prior, model and tool availability; each point retains ten scheduled sessions. All successes are in the aligned-prior stratum, where the correct law was supplied. **b**, Paired effects in five reused worlds and the equally weighted mean with its prespecified approximate 95% interval. Tool availability is the intervention: GPT used it in 0/30 sessions and DeepSeek in 20/30. Three failures remain in the denominator.

voluntary uptake and its message/computation costs. It does not isolate internal arithmetic, and the high versus medium settings do not imply equal reasoning budgets. The old B3 qualification fitter accessed a privileged simulator, so its success does not establish public-only identifiability. Historical schema differences were not randomized.

A post-hoc public-information inspection used all five existing worlds, one public packet per world and all eight calibration pairs per packet (40 pairs; 160 metric pairs; no missing rows). The candidate formulas name an effective coefficient and a process factor, but the packets omit complete observation equations, a process-factor definition and a noise model. Aqueous-product and phase-ratio observations are exactly equal between target and reference in 40/40 pairs each. Organic-product absolute differences have mean 0.005873 and maximum 0.012553; score differences have mean 0.004992 and maximum 0.010670. All paired values and per-world summaries are included in the supplement. These descriptive contrasts do not estimate signal-to-noise, test held-out prediction or prove that the implemented simulator is unidentifiable. This inspection adds no participant sessions, physical executions or independent worlds.

A.8 CROSS-MODEL SUCCESSOR DENOMINATORS

Historical stopped partials and stop boundaries remain immutable. They are not continued, pooled, or converted into successful cells. Each successor instead starts at its first scheduled unit with a complete failure-aware denominator.

The GPT-5.6-sol C2 successor scheduled all 135 task-world-prior cells and terminated with 126 completed, 3 failed, and 6 right-censored sessions, completing 1,253/1,260 participant experiments. Its provider-free evaluator completed 420/420 truth executions, scored 669/675 checkpoints and 129/135 laws, and launched 756/810 scheduled blind executions; 54 remained unstarted. The matched 135-cell DeepSeek/GPT analysis uses task-world cluster bootstrap intervals and never interprets model differences as provider effects.

The DeepSeek B3 successor similarly schedules 30 cells and retains 17 completed plus 13 participant-schema failures. Its historical canary is excluded rather than spliced into the new denominator. The GPT B3 cohort remains the independent completed 30-cell block described above.

The four-condition action successor schedules 45 strata by four conditions for each model: 180 slots per model and 360 total. DeepSeek reuses 45 immutable autonomous donors, of which 42 are eligible; GPT-5.6-sol creates a separate 45-donor cohort, of which 26 are eligible. No-evidence is scheduled for all strata, while yoked and learned-law recipients require the corresponding donor. Missing-donor descendants remain `not_started_due_to_missing_donor`. The primary analysis is the model-specific all-scheduled failure-aware strategy estimand over all 45 strata. DeepSeek/GPT autonomy-minus-no-evidence estimates are -0.0913 and +0.1102 with clustered intervals [-0.2124, 0.0388] and [-0.0533, 0.2794]. Donor-eligible estimates (-0.1214/-0.1379) are availability-conditioned sensitivities because eligibility is post-treatment; equal-task sensitivities are -0.0879 for DeepSeek and -0.0259 for GPT. The cross-model common donor-eligible population contains 26 strata from 13 task-world clusters. Yoked completion is 10/42 for DeepSeek and 24/26 for GPT, so autonomous-yoked contrasts mix scientific action quality with recipient-system failure and are not pure experiment-selection effects. The autonomous-minus-no-evidence task estimates were -0.2398/-0.4060/+0.3720 for DeepSeek and +0.0891/-0.2191/+0.4605 for GPT over electrochemistry/safety/crystallization, respectively. Thus the pooled direction masks strong task heterogeneity. All four registered contrasts and their bootstrap intervals are descriptive and were not multiplicity-adjusted. Conditions followed fixed manifest and dependency order rather than randomization or counterbalancing; DeepSeek autonomous donors also predated their recipients. Provider/time drift and condition-order effects are therefore not separable.

A.8.1 RECOVERY AND COMPOSITE PROVENANCE

Platform repair never changed a scientific question, coverage set, metric, or gate after observing an outcome. Affected blocks restarted from their first affected unit; discarded or superseded records remain in the failure timeline and do not enter the current estimate.

Table 5: Defect and recovery provenance for current publication surfaces.

Surface	Retained incident	Recovery boundary	Current inclusion
DeepSeek C2	one 15-cell shard affected by runtime/evaluator defects	whole shard restarted; evaluator rerun over the full composite	120 unaffected cells plus one 15-cell replacement shard
GPT C2	no inherited stopped cell	new 135-cell block from cell 1	all 135 scheduled outcomes
DeepSeek B3	historical three-cell canary failed schema	independent 30-cell successor from cell 1	17 complete plus 13 schema failures
Four-condition yoked	duplicate field deletion caused recipient <code>KeyError</code>	entire yoked condition restarted from its first admitted unit	recovered condition only; original failures retained in timeline
Oracle grids	path-length and evidence-reference limits	identical scientific block restarted after platform repair	only complete version-specific qualification blocks

The four-condition report remains a prospective development experiment rather than confirmatory evidence; the reanalysis does not rename it formal evidence. The DeepSeek C2 current composite is reported as a composite, not as one uninterrupted launch. Exact source bindings and excluded incident counts are in the supplement.

A.9 COMPLETE-PLAN ACTION ASSAY

In each open-action cell, a persistent agent performs twelve autonomous experiments and commits checkpoints after 0, 3, 6, 9, and 12 experiments. The host then reveals eight unseen ActionPlans. Each plan specifies ordered operations, all submitted parameters, initial-state assumptions, intermediate measurement positions, terminal assay, and declared omitted operations. Candidate outcomes, ranks, and other-arm evidence stay hidden. The evaluator verifies identity of disclosed, truth-evaluated, and replayed plans.

The primary endpoint is normalized regret of the selected plan within its task-world candidate set. Selected rank, Top-1, pairwise ordering, and law adequacy are secondary or diagnostic. Forty-five cells were scheduled across three tasks, five worlds, and three initial-model arms. Forty-two were eligible. The two agent-induced resource/process failures and one provider interruption are retained




Figure 8: Law quality and information strategies have distinct decision readouts. **a**, Law-implied versus participant regret in the 45-cell DeepSeek matrix; 42 observed choices are dots and three missing rankings are crosses above the plotting range. The diagonal marks equal regret. **b**, Failure-aware strategy means with valid/scheduled counts. **c**, Paired strategy-minus-no-evidence differences with task-stratified world-cluster bootstrap 95% intervals, retaining all 45 strata per model. Both autonomous contrasts cross zero. The strategy block is development evidence; yoked failures prevent a pure acquisition-effect interpretation. Complete-ranking diagnostics remain in the appendix.

as failures in the scheduled denominator and are not replaced. An independent repair is reported only as sensitivity evidence because it follows a new trajectory.

The continuous law–action analysis used only those same 42 originally eligible cells; the three failures stayed in the 45-cell scheduled denominator and were not imputed. Pearson and Spearman associations were computed pooled and by task. Uncertainty resampled the 15 frozen task–world clusters, retaining all eligible arms within a sampled cluster, with 10,000 bootstrap replicates and seed 20260827. The pooled Spearman coefficients were -0.073 for law MAE versus selected rank and -0.133 for law MAE versus normalized regret, with cluster-bootstrap 95% intervals $[-0.380, 0.256]$ and $[-0.452, 0.217]$. A prespecified sensitivity table swept law-MAE thresholds 0.05, 0.075, 0.10, 0.15, 0.20, 0.25, and 0.30 without selecting a cutoff from action outcomes.

A decision-aligned reanalysis used the last-available executable law from all 45 frozen DeepSeek-v4-flash longitudinal cells and reconstructed the same eight candidate feature packets without new participant, truth, or physics execution. Three cells had no terminal action ranking but retained an earlier executable law. For each law, truth-law error is the normalized regret of its implied Top-1. Action-utilization delta is participant regret minus law-implied regret. All 45 laws executed; none implied the true Top-1, whereas participant action selected Top-1 in 11/45 scheduled cells. Among 42 cells with a valid participant ranking, the participant followed the law-implied Top-1 in 12. Mean law-implied and participant regret were 0.438 and 0.344. This is a descriptive decomposition: neither law quality nor law use was randomized.

A.10 ORACLE-CONTROL QUALIFICATION

The unexecuted causal follow-up crossed three tasks, five worlds, three priors, and five conditions: no evidence, stepwise yoked evidence, autonomous exploration, learned law in a fresh context, and a provider-free oracle law. The planned denominator was 225 participant sessions, 45 autonomous donors, and 540 donor experiments. Participant execution required every fresh task–world oracle to pass candidate opportunity checks and to attain Spearman rank correlation at least 0.80 over the eight candidate plans using fit data disjoint from candidate outcomes.

The 96-query qualification stopped after eight of fifteen clusters. It completed 896/896 truth executions and exact replays, with eight candidate gates and seven oracle rank gates passing. The retained crystallization failure had $\rho = 0.738095$, Top-1 disagreement, and zero fit/candidate overlap. No operational canary or participant session began.

This failure initiated a transparent development sequence rather than a retrospective threshold change. Version 0.2 first failed on a fresh electrochemical world at $\rho = 0.785714$. Version 0.3 added exact typed-law distillation but failed on a fresh crystallization world at $\rho = 0.595238$. Version 0.4 fixed an ExtraTrees predictor: it replayed all 25 exposed construction worlds successfully, then stopped after five fresh units when the fifth reached only $\rho = 0.785714$ (four pass, one fail, ten unstated). Each version used a newly held-out fresh surface, stopped at its first registered failure, read no candidate outcomes during fitting, and made zero provider calls. These attempts are not pooled into a success rate and do not modify the original stop decision.

The 320-query construction held the fitted ExtraTrees family fixed while expanding coverage to 64 global and 256 candidate-neighborhood queries. A platform-defective partial caused by an evidence-reference schema limit was retained and not reused. After the contract repair, an independent construction run completed 2,352/2,352 truth executions and replays across seven exposed units, passing all seven and repairing four historical failures. The first prospective world then completed 336/336 executions and replays but failed at $\rho = 0.714286$; the remaining fourteen were not started. Candidate outcomes were never read during fitting.

The gate-alignment analysis added no execution. It re-read sixteen frozen unit versions and exactly reproduced their original Spearman and Top-1 values. Its action endpoints were Top-1, normalized regret, selection within 0.01 of optimum, and near-tie-aware pair ordering. It did not alter any historical result, threshold, or stop decision.

For eight candidates without ties, $\rho = 1 - \sum_i d_i^2/84$. Thus the retained 96-query failure at $\rho = 0.738095$ has $\sum_i d_i^2 = 22$, whereas the fresh 320-query failure at $\rho = 0.714286$ has 24. The scalar measures total displacement, not whether the best action moved.

Table 6: **Complete-ranking gates and action endpoints are different estimands.**

Frozen block	n	Rank pass	Top-1	≤ 0.01 regret
Fresh 96-query	8	7	1	3
Exposed 320-query	7	7	4	6
Fresh 320-query	1	0	1	1

A.11 INDEPENDENT-WORLD REPRESENTATION-BY-DECISION PROTOCOL

The factorial block fixes public evidence while replacing either its explicit representation or the decision rule. Each of two task families has five new public-test worlds. Two fixed model configurations each produce two independently sampled source artifacts per world, giving 40 source states nested within ten task–world clusters. Every source state has four conditions: model law with fresh agent selection (L-A), the same law with a deterministic maximizer (L-X), public-data fit with fresh agent selection (F-A), and the same fit with the maximizer (F-X). The scheduled design comprises 40 source and 80 recipient sessions, 160 condition slots, 120 public evidence experiments and 80 hidden candidate evaluations, with exact replay of each physical execution. This block uses fixed experiment acquisition.

Each world shares twelve evidence points and eight outcome-hidden candidate plans across models and repeats. Independently seeded Latin hypercubes fix the two-dimensional coordinates before

outcomes are observed. Electrochemistry varies controlled potential over 0.65–1.65 V and current over 20–120 mA. Its fixed probe uses 1.18 V, 70 mA and 630 s; controlled duration is 3540 s, reagent amount 0.004 mol, electrolyte profile 2 and solvent 1. The complete LHS is rejected before execution if any point violates minimum probe changes of 0.02 V or 1 mA. Crystallization varies reaction temperature over 350–405 K and crystallization temperature over 275–295 K. Catalyst 0 at 0.000315 mol, solvent 1, reagent 0.015 mol, stirring 675 rpm, reaction duration 3600 s, seed mass 0.008 g and crystallization duration 7200 s remain fixed. Both tasks share the same normalized design and disclose the complete executable plans.

L and F use the same unclipped quadratic basis $(1, x, y, x^2, xy, y^2)$ in linearly normalized coordinates. F is ridge regression on the twelve public utilities, with penalty 10^{-6} and an unpenalized intercept. Source models see evidence before candidates are revealed and return six finite coefficients. Fresh recipients receive evidence, candidates and one artifact, and return only a candidate identifier. Calls are tool-free, capped at 600 s each, with no repair turn or replacement. The request for 2,048 output tokens does not cap hidden reasoning; actual usage is recorded. Source states have a fixed outcome-blind permutation, and each world/model uses one L-first and one F-first decision repeat. The maximizer breaks exact ties by candidate order. All choices are sealed before candidate utilities are loaded for analysis.

Candidate utility is a single keyed-noise measured score in $[0, 1]$. M1 regret uses the fixed utility scale one; earlier action assays retain their registered range normalization. Missing or invalid choices receive regret one in the scheduled primary analysis; completed-only values retain separate denominators. Near-optimality means regret at most 0.01. The primary contrast is F-X minus L-X. Differences are first averaged over the four model/repeat states within world, then equally across worlds within task and across tasks. A task-stratified world bootstrap uses 20,000 percentile replicates. Its two-sided 95% upper endpoint must be below -0.01 to support the prespecified material benefit. The four secondary contrasts (L-X minus L-A, F-A minus L-A, F-X minus F-A, and their factorial interaction) use 98.75% marginal intervals, a Bonferroni adjustment for four comparisons. Five worlds per task give approximate small-sample intervals; model repeats do not add independent worlds.

The nearest-public-observation baseline uses normalized Euclidean distance and row-order ties; the random baseline is its exact uniform-candidate expected regret. Both reuse the same data and count once per world. Missing source artifacts block the dependent L conditions while F continues. Physical, semantic or replay failure stops dependent provider work; forbidden tools and missing or reused session identity invalidate the remaining block. Interrupted started slots are retained without retries. Physical CPU/wall includes replay; recipe-resource totals count primary execution once. The protocol replaces explicit artifacts and decision computation, and does not identify internal belief mediation or artifact-only transfer: recipients still receive the original evidence.

A.11.1 FACTORIAL OUTCOMES AND RESOURCES

All 40 source states, 120 sessions and 160 conditions completed; there were no failed, blocked, replaced or unstarted slots. Scheduled and completed-only losses coincide. Mean regrets in L-A/L-X/F-A/F-X order were 0.01436/0.01502/0.00425/0.00425 for DeepSeek and 0.00425/0.00425/0.00425/0.00425 for GPT. These configuration-specific summaries are descriptive. F-X selected the exact best measured candidate in 20/40 states and a candidate within 0.01 in 28/40; the replicated fit copies still represent only ten worlds. Candidate MAE was 0.05028 for DeepSeek laws, 0.04832 for GPT laws and 0.04459 for the public fit, all finite. Better average prediction error thus did not establish a material decision improvement.

The ten-world average nearest-evidence regret was 0.00354, versus 0.00425 for F-X and 0.11109 for exact expected uniform-random choice. These prespecified baselines reuse the same evidence and candidate evaluations; the comparison supplies no novel-algorithm advantage. The F-A/F-X bootstrap interval collapses to zero because all observed paired choices agree. It does not establish exact equivalence for unobserved worlds or model samples.

Physical execution and replay consumed 1,780.8 s wall time and 1,768.3 s CPU. The 120 public evidence executions account for 1,068.6 s wall time; 80 hidden evaluations account for 712.3 s. Recipe resources count primary executions once: 2,300 operations, 54.8 measurement units, 1,497,000 simulated recipe seconds and 1.9 mol reagent. Provider calls used 8,859.5 s wall time,




Figure 9: **Fixed-evidence intervention did not establish a material repair benefit.** **a**, Every world has its own row. Hollow and filled points show model-law and fitted-law regret under the same maximizer, averaged over two models and two repeats; concentric points denote equality. Blue denotes electrochemistry and purple crystallization. **b**, Colored points are individual world effects; black diamonds and lines show means and prespecified intervals (95% primary; 98.75% secondary). The dotted primary reference is -0.01. Interaction is $(R_{F-X} - R_{L-X}) - (R_{F-A} - R_{L-A})$ in regret. Zero-width intervals describe observed equality, not population equivalence. All 160 selections were available; repeats do not increase the ten-world denominator.

Table 7: **Prespecified factorial contrasts.** Differences use the fixed regret scale. The primary material threshold is -0.01. Secondary intervals adjust for four comparisons.

Contrast	Mean difference	Interval
F-X minus L-X (primary)	-0.005384	95% [-0.016303, 0.000614]
L-X minus L-A	+0.000331	98.75% [0.000000, 0.001323]
F-A minus L-A	-0.005053	98.75% [-0.017329, 0.001582]
F-X minus F-A	0.000000	98.75% [0.000000, 0.000000]
Interaction	-0.000331	98.75% [-0.001323, 0.000000]

1,623,791 input tokens (450,816 cached) and 942,258 output tokens (939,441 reasoning). All 120 calls report input/output usage; reasoning is included in output and cache in input. Source generation used 7,404.1 s versus 1,455.4 s for decisions. These costs describe this block, without extrapolating provider billing or virtual resource use to laboratory expenditure.

A.12 INFORMATION SEPARATION AND SAME-WORLD CONTEXT PORTABILITY

The context-portability assay reuses all ten factorial worlds, their twelve public observations per world and forty sealed source states. Private mechanisms, initial-state construction, objective, control ranges and fixed controls remain unchanged. A separately fixed eight-point Latin-hypercube design provides new candidate plans, disjoint from both the source observations and the earlier candidates. This requires 80 hidden physical evaluations and 80 exact replays, but no additional source measurements and no measurements by recipients. The independent-world denominator remains ten; the forty sources and 160 recipient sessions are nested within them. This is a new-decision assay in the same worlds, not transfer across changed private parameters or mechanism topology.

Four information conditions receive identical task descriptions, coordinate/basis semantics, utility direction and range, complete candidate controls and ActionPlans, and the same minimal candidate-ID output schema. None receives no experiment-derived information. Raw receives the twelve public observations, including their plans, without a law. L and F receive only the six coefficients of the sealed model-generated or fitted quadratic, respectively, without raw observations, source dialogue or provenance labels. Coefficients are reused exactly, without refitting to the new candidates. No condition receives prior candidate outcomes or world identity. Information quantity and context length are part of these deployment strategies; prompts are not padded to equal length. Prompt bytes and actual provider tokens are recorded separately.

The same two model configurations and two repeats per world are used, with fresh tool-free recipients. The forty source states are permuted before execution, and cyclic rotations of the four conditions across the nested states place each condition in each serial position once per world. Sessions have a 600-second timeout and no repair turns or retries. Missing artifacts block only their dependent information condition. Participant schema failures remain in the scheduled denominator; interrupted started calls are retained without replacement. Physical, replay, information-isolation or session-identity defects stop dependent work and invalidate formal inference. All recipient and deterministic choices are sealed before analysis loads new scores.

The primary contrast is L minus none in failure-aware regret. The utility scale is fixed at one; missing or invalid choices receive regret one. Within-world differences average the four nested model/repeat states, then weight worlds equally within task and tasks equally overall. A 20,000-replicate task-stratified world bootstrap supplies a two-sided 95% interval; its upper endpoint must be below -0.01 to support a material benefit. Five secondary contrasts—raw minus none, F minus none, L minus raw, F minus raw and F minus L—use 99% marginal intervals, adjusting for five comparisons. Completed-only losses retain eligible denominators. Top-1, regret at most 0.01, individual-world effects and exact-maximizer agreement accompany the primary estimate. Small-sample and previously exposed-world limits remain; nonsignificance is not equivalence.

Deterministic L/F maximizers, nearest public evidence and uniform-random expected regret are prespecified descriptive controls on the same new candidates. They add no provider or physical executions. New recipient and hidden-evaluation costs are separate from historical acquisition and source-generation costs. None uses no acquired evidence; raw and F inherit the shared public acquisition, while L also inherits its model’s original source-generation cost. Source tool permissions remain those of the factorial experiment: matching recipient permissions does not remove the original model-versus-numerical-solver arithmetic difference. Zero recipient queries do not estimate experimental savings, and no internal mediation or new repair algorithm is identified by this information comparison.

A.12.1 INFORMATION-SEPARATION OUTCOMES AND RESOURCES

All 160 recipients completed with no failures, retries or replacements; scheduled and completed-only losses coincide. The primary effect was negative in nine worlds and zero in one. Benefits were substantially larger in electrochemistry than crystallization.

Table 8: Information-separated decision quality. Every model/condition has 20 scheduled and completed recipients, nested within ten reused worlds.

Model	Information	Mean regret	Near-optimal	Top-1
deepseek	none	0.15270	3/20	3/20
deepseek	raw	0.01643	17/20	15/20
deepseek	L	0.01157	11/20	9/20
deepseek	F	0.02309	12/20	10/20
gpt	none	0.14184	10/20	10/20
gpt	raw	0.00605	13/20	11/20
gpt	L	0.00850	12/20	11/20
gpt	F	0.00609	12/20	10/20

Model-law recipients matched the exact maximizer in 40/40 states and fitted-law recipients in 39/40. Deterministic model-law and fitted-law regret was 0.01004 and 0.00609. The single fitted-law

Table 9: Prespecified information contrasts. Negative regret differences favor the first condition. The primary material threshold is -0.01.

Contrast	Mean difference	Interval
L minus none	-0.137231	95% [-0.155845, -0.122566]
raw minus none	-0.136032	99% [-0.166955, -0.106025]
F minus none	-0.132680	99% [-0.164488, -0.099817]
L minus raw	-0.001199	99% [-0.023529, 0.009508]
F minus raw	+0.003352	99% [-0.021421, 0.037629]
F minus L	+0.004551	99% [-0.008013, 0.031772]

recipient disagreement increases its group mean to 0.01459; it is retained as a valid choice. Nearest public evidence selected the measured optimum in all ten worlds; uniform-random expected regret was 0.12022. These comparisons do not establish a novel-algorithm advantage.

New physical execution/replay used 724.6 s wall time and 718.3 s CPU, with 920 operations, 21.92 measurement units, 598,800 simulated recipe seconds and 0.76 mol reagent counted once for primary executions. Provider time was 2,722.6 s, with usage reported for 160/160 calls: 1,716,721 input tokens (648,704 cached) and 255,544 output tokens (253,304 reasoning). Raw versus model-law deployment used 541,225/392,101 input tokens and 102,235/47,745 output tokens. Prompt bytes were 838,900/351,996 (58.0% fewer), whereas provider input fell by 27.6%; these measures include different overheads and must not be interchanged. Raw/model-law provider time was 929.1/529.3 s. Costs describe this deployment block, not monetary expenditure or experimental savings. Historical M1 public acquisition and original source generation are reported separately and were not incurred again. No equivalence or new-mechanism transfer claim follows from these data.

A.13 OBSERVATION-MAPPING DISCLOSURE INTERVENTION

This block uses ten new partition worlds with independently fixed seeds and target exponents spanning 1.35–2.25 within the same power-law family. Three prior conditions cross original formula disclosure with the complete observation mapping in each world and configuration. Each model has 60 scheduled fresh sessions, with a pre-response and a post-evidence continuation in the same thread. All 119 delivered DeepSeek prompts match their GPT counterparts byte for byte. Both conditions expose identical candidate parameter domains, observations, queries, prior descriptions, schema and calculator access. The post-response in the complete condition additionally receives the observation equations, process factors and noise specification. Disclosure length is consequently bundled with information completeness. Observation noise is paired across conditions; no unseen-family or changed-mechanism generalization is tested.

Before provider execution, a public-only numerical reference was qualified on the current observation contract and all ten stimulus worlds. The first 24-execution endpoint block exposed an inactive runtime setting and failed qualification; its executions and replays remain recorded. After correcting the new block’s setting, all 24 requalification and 240 ten-world stimulus executions and their exact replays passed. This correction was made before the agent blocks; it does not change the historical B3 or optional-numeric data. Six development sessions per configuration passed separately before the 60-session blocks.

GPT uses medium reasoning and DeepSeek Flash uses low reasoning in the same native agent harness. Each turn is limited to 600 seconds and each session to 1,200 seconds, with at most eight post-response calculator attempts and zero provider retries. The initial GPT block limit was four hours. DeepSeek’s original calendar limit was later removed by user instruction; after 21 terminal sessions, the remaining 39 were dispatched by one coordinator to at most three independent workers, finishing in 99.44 minutes without a rate-limit downgrade. These are disclosed scheduling changes. Scientific inputs, per-session budgets and failure scoring remained fixed; the systems are not equal-compute competitors. A thinking-disabled development configuration failed qualification and did not enter either formal result.

Family must be correct and exponent absolute error at most 0.10 for joint recovery. The primary population contains the two unknown/wrong-prior strata; correct-prior retention is reported separately.

Failures receive zero joint recovery and Top-1, and worst normalized regret of one. Prediction error is reported only for available valid submissions with exact sample counts. The complete-minus-original recovery contrast averages priors and configurations within world, then weights the ten worlds equally. Approximate 95% intervals resample these ten clusters 20,000 times with a fixed analysis seed; model/query counts do not increase independent worlds.

Table 10: Original disclosure results. Unknown/wrong-prior groups each contain 20 scheduled sessions. Parentheses in the prediction column give available sample counts. Correct-prior retention is separate.

Configuration	Disclosure	Recovery	Top-1	MAE (n)	Regret
GPT / medium	Original	6/20	5/20	0.03871 (20)	0.07097
	Complete	20/20	12/20	0.00606 (20)	0.00748
DeepSeek / low	Original	3/20	6/20	0.04654 (18)	0.15090
	Complete	17/20	4/20	0.00608 (17)	0.24263

GPT completes 60/60 sessions; DeepSeek completes 52/60, with two tool-budget failures, three post-response schema failures, one provider output-limit failure, one turn timeout and one retained user interruption. Together they account for 239/240 scheduled turns. In the complete DeepSeek recovery group, all 17 valid submissions recover structure, but the primary remains 17/20. Correct-prior retention changes from 8/10 to 10/10 for GPT and 6/10 to 10/10 for DeepSeek. Recovery improves in eight GPT worlds (two ties) and all ten DeepSeek worlds. Both effects are +70 percentage points; configuration-specific approximate intervals are [45, 95] and [55, 85]. The pooled recovery rate changes from 9/40 to 37/40, with interval [52.5, 85] for the gain. The opposite action directions are descriptive and do not establish significant worsening or an internal structure-to-action mechanism.

Resources are cumulative within a conversation, so the last available usage per thread is counted once. GPT reports 1,552,862 input tokens (332,544 cached), 327,180 output tokens (278,923 reasoning), and 8,355.718 session-wall seconds. DeepSeek reports 15,067,287 input tokens (13,120,384 cached), 3,470,258 output tokens (3,356,264 reasoning), 296 calculator attempts and 88 rejections. DeepSeek session-wall time totals 26,198.564 seconds, with the interrupted session incomplete; it is not parallel elapsed time. Usage is available for 114/119 attempted turns, so DeepSeek resource totals are lower bounds. No cost is erased by concurrency, interruption or failure.

A.13.1 ONE ADDITIONAL ATTEMPT FOR EACH FAILED SESSION

After inspecting the original terminal results, the user authorized one additional attempt for every failed or interrupted DeepSeek session. This selected-failure sensitivity block fixed all eight eligible cases before execution, retained the original scientific inputs and frozen executor, and used a fresh complete pre/post session for each. Reasoning, turn/session limits, calculator limits and zero provider retries were unchanged. Three independent workers were allowed, without a block calendar deadline. Further failures or unfavorable answers would not trigger another attempt. The original 120-session primary and its eight failures remain intact.

All eight additional sessions completed validly (16/16 turns), without execution failures. Five unknown/wrong-prior cases recovered structure; two of three correct-prior cases retained it, while the remaining valid answer was incorrect. In the descriptive virtual result allowing one retry per original failure, DeepSeek recovery changes from 5/20 to 20/20, compared with 3/20 to 17/20 in the original primary. Top-1 is 7/20 versus 5/20 and normalized regret is 0.05551 versus 0.09597. Thus the structure benefit and the descriptive adverse action direction persist after this reliability sensitivity. Correct-prior retention is 8/10 versus 10/10. These are post-hoc selected attempts, not an independent replication, a new primary or evidence of statistically significant action worsening. No confidence interval is claimed for this selected retry comparison.

Additional elapsed time was 20.10 minutes; session-wall time summed to 3,404.030 seconds. All 16 usage receipts are available: 2,255,824 input tokens (1,979,776 cached), 450,242 output tokens (433,005 reasoning), 35 calculator attempts and 11 rejections. These costs are additional to the original formal block. The complete programme therefore retains 128 attempted sessions, including the original eight failures, rather than erasing them after successful retries.

1458
1459
1460
1461
1462
1463
1464
1465
1466
1467
1468
1469
1470
1471
1472
1473
1474
1475
1476
1477
1478
1479
1480
1481
1482
1483
1484
1485
1486
1487
1488
1489
1490
1491
1492
1493
1494
1495
1496
1497
1498
1499
1500
1501
1502
1503
1504
1505
1506
1507
1508
1509
1510
1511

Table 11: **Descriptive one-retry sensitivity for DeepSeek.** Original successful attempts are retained; each original failure contributes its single additional attempt. The original primary is unchanged.

Population	Disclosure	Recovery/retention	Top-1	Regret
Unknown/wrong	Original	5/20	7/20	0.05551
	Complete	20/20	5/20	0.09597
Correct prior	Original	8/10	5/10	0.06583
	Complete	10/10	3/10	0.08297

A.14 ADDITIONAL SCOPE BOUNDARIES

The study covers simulated chemistry and two model families in study-specific configurations, not laboratory fidelity or universal chemical coverage. C2 and B3 have matched scheduled surfaces; A-P/B2 have complete separate denominators plus a B2-only DeepSeek-low ablation. Differential failures and non-random provider assignment prohibit model ordering. The public cohort is not private confirmation, and transfer across changed physical mechanisms remains untested. The autonomous open-action cohort lacks a same-agent no-evidence baseline; the oracle-free successor adds four information strategies, but donor/recipient failures limit mediation claims. The original five-condition oracle cohort remains unexecuted.