



Initial models redirect search, but evidence does not selectively repair the wrong model

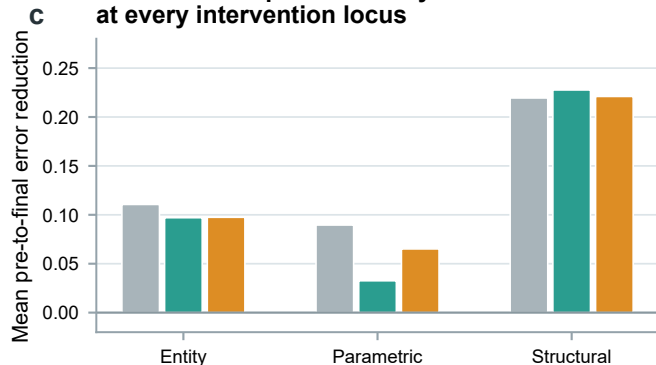
Starting predictions are arm dependent, not uniformly ordered by correctness



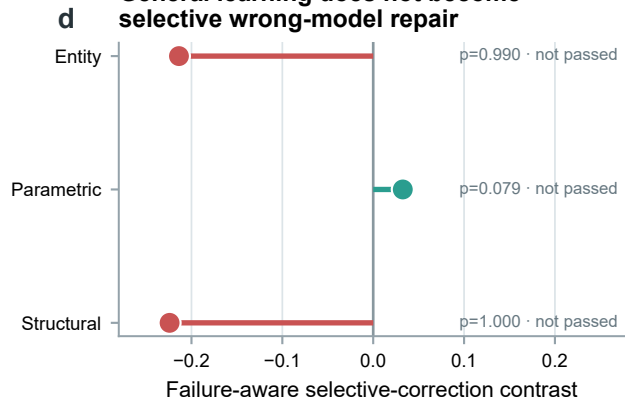
The intervention enters the first experimental trajectory



Predictions improve in every arm at every intervention locus



General learning does not become selective wrong-model repair



Prospective cohort: 45 matched task-world clusters and 135 persistent campaigns. Panels a–c are descriptive manipulation and learning summaries; first-recipe divergence has no same-arm replicate baseline. Panel d reports the prespecified failure-aware locus decisions; positive values favor selective repair.