





Matched evidence yields numerical convergence without reliable law-to-action recovery

45/45 structural sessions across 3 configurations

B3: 5/60 correct laws; 0 useful action gains

● Opaque ● Aligned ● Misindexed



Panels a-b show the corrected DeepSeek-high B2 structural assay; panel c adds matched GPT-medium and DeepSeek-low evidence. All three B2 configurations put misspecified cells near the observed outcomes yet recover the exact 1.75 law in 0/5 worlds. Panel d reports the separate participant-identifiable B3 surface: DeepSeek has 13 retained schema failures, Codex recovers 5/30 joint laws, and neither model realizes a registered useful action gain. Five-world and model comparisons are descriptive; no capability-ranking test is performed.