


Evidence, executable laws and evaluator rankings are imperfect action proxies

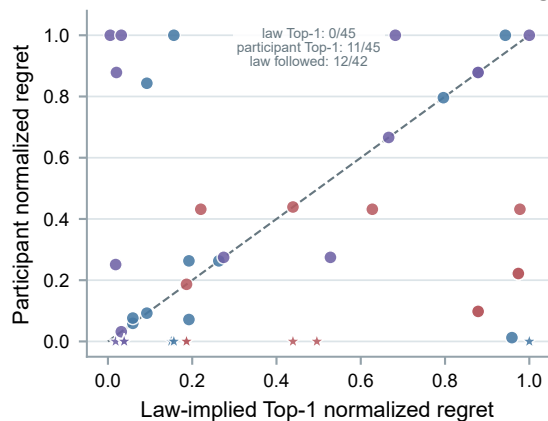
● DeepSeek-v4-flash high ◆ GPT-5.6-sol medium ★ Top-1 selected

● Electro ● Crystal ● Safety

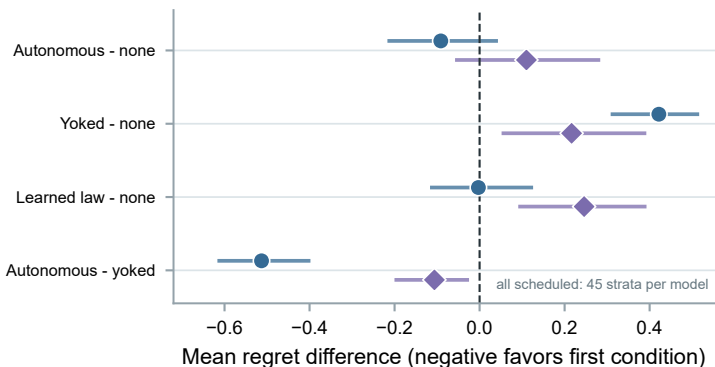
a All-scheduled action outcomes differ by model and information strategy



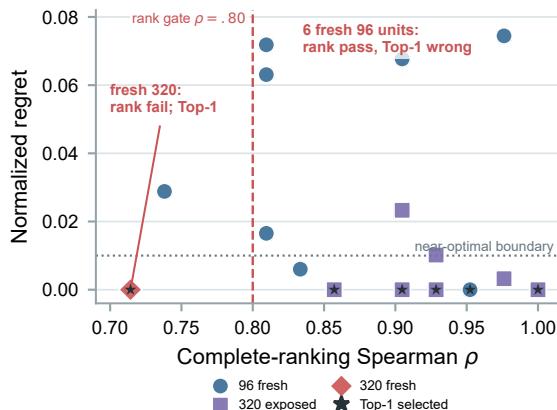
b DeepSeek actions often depart from the executable law's candidate ordering



c Strategy estimates differ by model; no consistent autonomy benefit is established



d Full-ranking validity and action validity disagree in both directions



Four-condition extension: 360 scheduled slots and 45 strata per model; donor, blocked-recipient and recipient failures remain in the primary strategy estimates. Open-action cohort: 45 scheduled cells and 240/240 truth/replay. Oracle diagnostics contain zero participant sessions and 16 frozen unit-version records. Full-ranking correlation is secondary; model differences are not provider causal effects.